


Decoder-Only Transformer — Model Architecture



defaults: $d_{\text{model}}=512 \cdot N=30$ layers · 8 Q heads / 2 KV heads · tied embeddings